

M3. Actividad 4. Modelos de regresión lineal y sistemas de recomendaciones


```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from sklearn.model_selection import train_test_split
import statsmodels.stats as sms
import statsmodels.api as sm
```

```
In [ ]: data = pd.read_csv(r'C:\Users\Raul\OneDrive\Escritorio\CS\TC2004B.101\M3-Act4\menu.
data.head()
```

Out[]:

	Category	Item	Serving Size	Calories	Calories from Fat	Total Fat	Total Fat (% Daily Value)	Saturated Fat	Saturated Fat (% Daily Value)	Tr
0	Breakfast	Egg McMuffin	4.8 oz (136 g)	300	120	13.0	20	5.0	25	
1	Breakfast	Egg White Delight	4.8 oz (135 g)	250	70	8.0	12	3.0	15	
2	Breakfast	Sausage McMuffin	3.9 oz (111 g)	370	200	23.0	35	8.0	42	
3	Breakfast	Sausage McMuffin with Egg	5.7 oz (161 g)	450	250	28.0	43	10.0	52	
4	Breakfast	Sausage McMuffin with Egg Whites	5.7 oz (161 g)	400	210	23.0	35	8.0	42	

5 rows × 24 columns



```
In [ ]: data.isna().any()
```

```
Out[ ]: Category          False
        Item              False
        Serving Size      False
        Calories           False
        Calories from Fat  False
        Total Fat          False
        Total Fat (% Daily Value) False
        Saturated Fat      False
        Saturated Fat (% Daily Value) False
        Trans Fat          False
        Cholesterol         False
        Cholesterol (% Daily Value) False
        Sodium             False
        Sodium (% Daily Value) False
        Carbohydrates       False
        Carbohydrates (% Daily Value) False
        Dietary Fiber       False
        Dietary Fiber (% Daily Value) False
        Sugars              False
        Protein             False
        Vitamin A (% Daily Value) False
        Vitamin C (% Daily Value) False
        Calcium (% Daily Value) False
        Iron (% Daily Value) False
        dtype: bool
```


```
In [ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 260 entries, 0 to 259
Data columns (total 24 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Category                                   260 non-null    object
1   Item                                       260 non-null    object
2   Serving Size                             260 non-null    object
3   Calories                                  260 non-null    int64
4   Calories from Fat                        260 non-null    int64
5   Total Fat                               260 non-null    float64
6   Total Fat (% Daily Value)               260 non-null    int64
7   Saturated Fat                           260 non-null    float64
8   Saturated Fat (% Daily Value)           260 non-null    int64
9   Trans Fat                               260 non-null    float64
10  Cholesterol                              260 non-null    int64
11  Cholesterol (% Daily Value)             260 non-null    int64
12  Sodium                                   260 non-null    int64
13  Sodium (% Daily Value)                  260 non-null    int64
14  Carbohydrates                           260 non-null    int64
15  Carbohydrates (% Daily Value)           260 non-null    int64
16  Dietary Fiber                           260 non-null    int64
17  Dietary Fiber (% Daily Value)           260 non-null    int64
18  Sugars                                   260 non-null    int64
19  Protein                                  260 non-null    int64
20  Vitamin A (% Daily Value)               260 non-null    int64
21  Vitamin C (% Daily Value)               260 non-null    int64
22  Calcium (% Daily Value)                 260 non-null    int64
23  Iron (% Daily Value)                    260 non-null    int64
dtypes: float64(3), int64(18), object(3)
memory usage: 48.9+ KB
```

```
In [ ]: # data = pd.get_dummies(data, drop_first=True)
variables = ['Total Fat', 'Carbohydrates', 'Protein', 'Sodium', 'Calories']
data = data[variables]
data.head()

fig = plt.figure(figsize=(15,3))
counter = 0
for column in data:
    counter += 1
    plt.subplot(1,5, counter)
    sns.histplot(data, x=column)

plt.tight_layout()
```



```
In [ ]: X = data.drop('Calories', axis=1)
        y = data['Calories']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_s
```

0. Funciones

```
In [ ]: def forward_selection(X_train, y_train, model, metric):
        testing_features = list(X_train.columns)
        final = []
        scores = [(-1, None)]
        highscore = -1

        while (max(scores)[0] < metric):
            if (len(testing_features) == 0):
                print('Metric not met.')
                print('R2:', highscore)
                return final

            scores = []
            for column in X_train[testing_features]:
                features = final + [column]
                X_temp = X_train[features]
                model.fit(X_temp, y_train)
                score = model.score(X_temp, y_train)
                scores.append((score, column))

            if (highscore < max(scores)[0]):
                final.append(max(scores)[1])
                highscore = max(scores)[0]

            testing_features.remove(max(scores)[1])
            print('Depth:', len(final), ' - ', final, highscore)
        print(final)
        return final
```

1. Lineal Multivariante

```
In [ ]: from sklearn.linear_model import LinearRegression
        reg1 = LinearRegression()

        features = forward_selection(X_train, y_train, reg1, 0.99)
        X_train_1 = X_train[features]
        X_test_1 = X_test[features]


        reg1.fit(X_train_1, y_train)

        y_pred_1 = reg1.predict(X_test_1)

        plt.figure(figsize=(12, 4))
        sns.scatterplot(x=y_test.index, y=y_test, color='blue', label='Actual')
        sns.scatterplot(x=y_test.index, y=y_pred_1, color='red', label='Predicted')
```

```
plt.xlabel('Index')
plt.ylabel('Target')
plt.title('Actual vs Predicted')
plt.legend()
plt.show()
```

Depth: 1 - ['Total Fat'] 0.8132128731140438
 Depth: 2 - ['Total Fat', 'Carbohydrates'] 0.9870735117068954
 Depth: 3 - ['Total Fat', 'Carbohydrates', 'Protein'] 0.999480255395087
 ['Total Fat', 'Carbohydrates', 'Protein']



```
In [ ]: from sklearn.metrics import r2_score

n = len(X_train_1) # Number of registers
k = len(X_train_1.columns) # Number of columns

r2 = r2_score(y_pred_1, y_test)
r2_adj = 1 - (1-r2)*(n-1)/(n-k-1)

col1 = pd.DataFrame({'Lineal Multivariante' : [r2, r2_adj]}, index=['R^2', 'R^2 Adj'])
```

2. Polinomial Multivariante

```
In [ ]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

poly_features = PolynomialFeatures(degree=3, interaction_only=False, include_bias=False)
X_poly_train = pd.DataFrame(poly_features.fit_transform(X_train))
X_poly_test = pd.DataFrame(poly_features.transform(X_test))

reg2 = LinearRegression()
features = forward_selection(X_poly_train, y_train, reg2, 0.99)
X_train_2 = X_poly_train[features]
X_test_2 = X_poly_test[features]


# Ajustar el modelo a los datos
reg2.fit(X_train_2, y_train)

y_pred_2 = reg2.predict(X_test_2)

plt.figure(figsize=(12, 4))
```

```
sns.scatterplot(x=y_test.index, y=y_test, color='blue', label='Actual')
sns.scatterplot(x=y_test.index, y=y_pred_2, color='red', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Target')
plt.title('Actual vs Predicted')
plt.legend()
plt.show()
```

```
Depth: 1 - [5] 0.8175422978281982
Depth: 2 - [5, 1] 0.8794855463590275
Depth: 3 - [5, 1, 0] 0.9872538558739244
Depth: 4 - [5, 1, 0, 2] 0.9994804759036104
[5, 1, 0, 2]
```



```
In [ ]: from sklearn.metrics import r2_score

n = len(X_train_2) # Number of registers
k = len(X_train_2.columns) # Number of columns

r2 = r2_score(y_pred_2, y_test)
r2_adj = 1 - (1-r2)*(n-1)/(n-k-1)

col2 = pd.DataFrame({'Polinomial Multivariante' : [r2, r2_adj]}, index=['R^2', 'R^2
```

3. Polinomial Multivariante con Interacciones entre las Variables de Entrada

```
In [ ]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

poly_features = PolynomialFeatures(degree=3, interaction_only=True, include_bias=False)
X_poly_train = pd.DataFrame(poly_features.fit_transform(X_train))
X_poly_test = pd.DataFrame(poly_features.transform(X_test))

reg3 = LinearRegression()
features = forward_selection(X_poly_train, y_train, reg3, 0.99)
X_train_3 = X_poly_train[features]
X_test_3 = X_poly_test[features]

# Ajustar el modelo a los datos
```

```

reg3.fit(X_train_3, y_train)

y_pred_3 = reg3.predict(X_test_3)


plt.figure(figsize=(12, 4))
sns.scatterplot(x=y_test.index, y=y_test, color='blue', label='Actual')
sns.scatterplot(x=y_test.index, y=y_pred_3, color='red', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Target')
plt.title('Actual vs Predicted')
plt.legend()
plt.show()

```

```

Depth: 1 - [4] 0.8175422978281982
Depth: 2 - [4, 1] 0.8794855463590275
Depth: 3 - [4, 1, 0] 0.9872538558739244
Depth: 4 - [4, 1, 0, 2] 0.9994804759036104
[4, 1, 0, 2]

```



```

In [ ]: from sklearn.metrics import r2_score

n = len(X_train_3) # Number of registers
k = len(X_train_3.columns) # Number of columns

r2 = r2_score(y_pred_3, y_test)
r2_adj = 1 - (1-r2)*(n-1)/(n-k-1)

col3 = pd.DataFrame({'Polinomial Multivariante con Interacciones' : [r2, r2_adj]},

```

4. Evaluación de Métricas

```

In [ ]: metrics = pd.concat([col1, col2, col3], axis=1)
metrics

```

Out[]:

	Lineal Multivariante	Polinomial Multivariante	Polinomial Multivariante con Interacciones
R²	0.999442	0.999441	0.999441
R² Adjusted	0.999433	0.999430	0.999430

5. Supuestos de Regresión Lineal

```
In [ ]: def vibe_check(y_test, y_preds):
    fig = plt.figure(figsize=(16,16))
    fig.suptitle('Residual Analysis', fontsize=20) # Add title here
    index = 0

    for y_pred in y_preds:
        residuals = y_test - y_pred

        # Plot residuals vs fitted values to check for independence
        plt.subplot(3,4,1 + 4*index)
        plt.scatter(y_pred, residuals)
        plt.axhline(y=np.mean(residuals), color='b', linestyle='-')
        plt.title('Residuals vs Fitted Values')
        plt.xlabel('Fitted Values')
        plt.ylabel('Residuals')
        plt.grid(True)

        # Plot histogram of residuals
        plt.subplot(3,4,2 + 4*index)
        sns.histplot(residuals, kde=True, bins=10)
        plt.title('Histogram of Residuals')
        plt.xlabel('Residuals')
        plt.ylabel('Frequency')
        plt.grid(True)

        # Plot QQ plot of residuals
        plt.subplot(3,4,3 + 4*index)
        stats.probplot(residuals, dist="norm", plot=plt)
        plt.title('QQ Plot of Residuals')
        plt.grid(True)

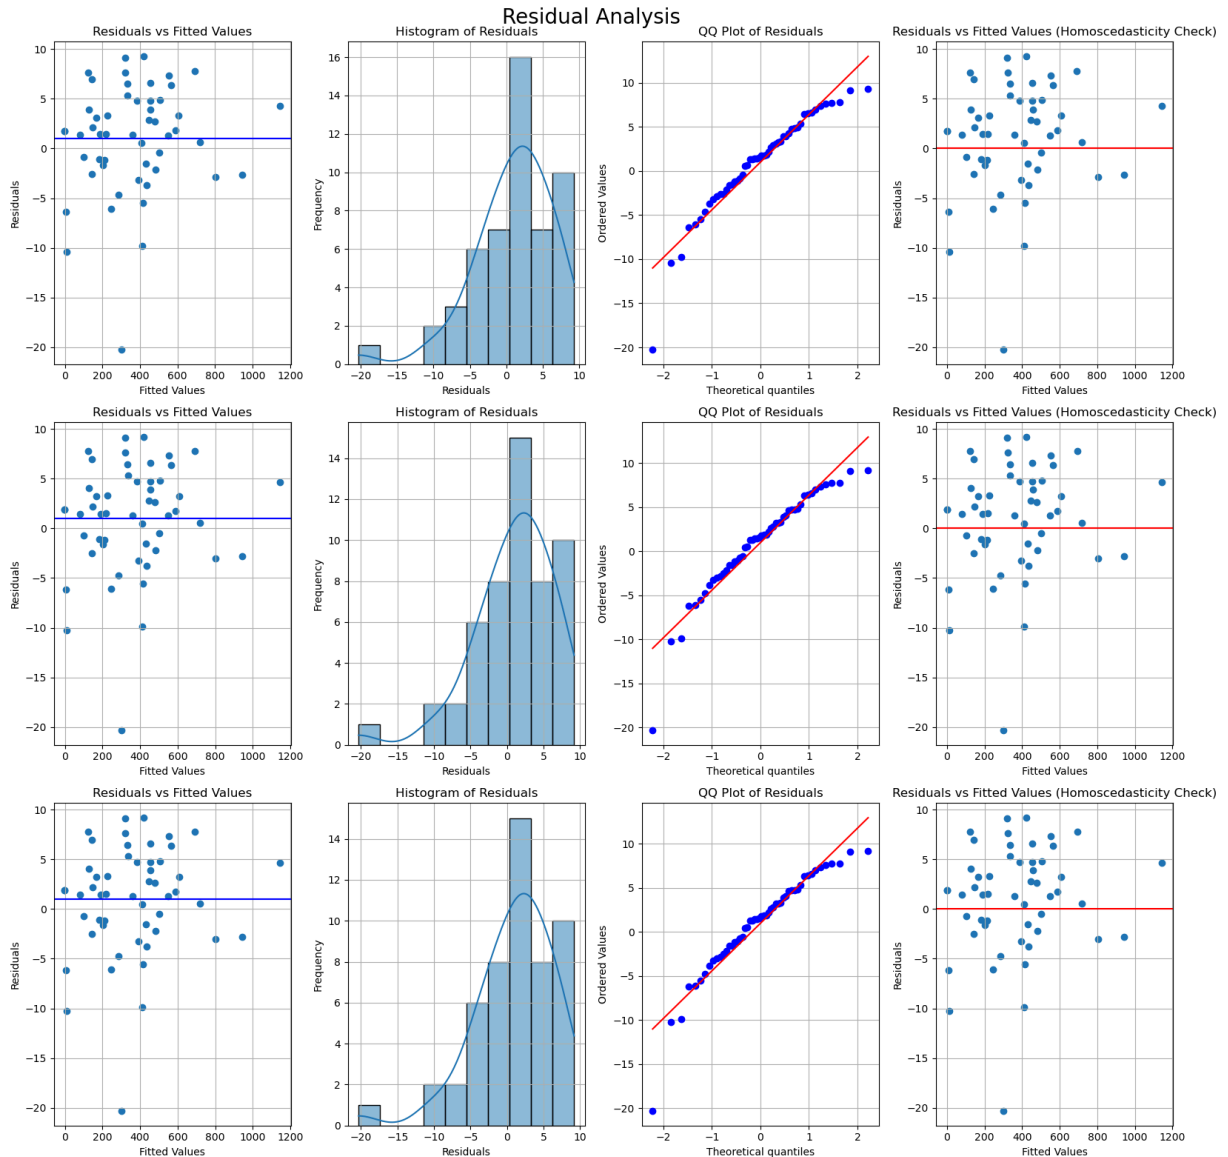
        # Check homoscedasticity using predicted values and residuals
        plt.subplot(3,4,4 + 4*index)
        plt.scatter(y_pred, residuals)
        plt.axhline(y=0, color='r', linestyle='-')
        plt.title('Residuals vs Fitted Values (Homoscedasticity Check)')
        plt.xlabel('Fitted Values')
        plt.ylabel('Residuals')
        plt.grid(True)

        index += 1
```



```
plt.tight_layout()
plt.show()
```

```
In [ ]: vibe_check(y_test, [y_pred_1, y_pred_2, y_pred_3])
```



Supuestos de una Regresión Lineal

1. **Linealidad:** Se dice que un modelo de regresión lineal cumple el supuesto de linealidad cuando la ecuación utilizada representa adecuadamente una relación lineal entre variables dependientes e independientes.
2. **Normalidad:** Se dice que un modelo de regresión lineal cumple el supuesto de normalidad cuando los residuos del modelo siguen una distribución normal. Para verificar este supuesto, se puede utilizar la prueba de Shapiro-Wilk.
3. **Homocedasticidad:**

Se dice que un modelo de regresión lineal cumple el supuesto de homocedasticidad cuando la varianza de los residuos es constante a lo largo de los valores de las variables

independientes. Para verificar este supuesto, se puede utilizar la prueba de Breusch-Pagan.

4. **Independencia:** Se dice que un modelo de regresión lineal cumple el supuesto de independencia cuando los residuos del modelo no están correlacionados entre sí. Para verificar este supuesto, se puede utilizar la prueba de Durbin-Watson.

```
In [ ]: from statsmodels.stats.diagnostic import het_breuschpagan, normal_ad
from statsmodels.stats.stattools import durbin_watson
from statsmodels.tools.tools import add_constant

def check_assumptions(y_test, y_preds):
    index = 1
    for y_pred in y_preds:
        residuals = y_test - y_pred

        # Normality test
        test_statistic, p_value_normality = normal_ad(residuals)
        print(f"Model {index}")
        print(f"Normality Test (Shapiro-Wilk): p-value = {p_value_normality:.4f}")

        # Homoscedasticity test
        # Add a constant column to X_test
        X_test_with_const = add_constant(X_test)
        _, p_value_homoscedasticity, _, _ = het_breuschpagan(residuals, X_test_with_const)
        print(f"Homoscedasticity Test (Breusch-Pagan): p-value = {p_value_homoscedasticity:.4f}")

        # Independence test
        durbin_watson_statistic = durbin_watson(residuals)
        print(f"Durbin-Watson Statistic: {durbin_watson_statistic:.4f}")

        index += 1
        print("-----")

    # Check assumptions for each model
    check_assumptions(y_test, [y_pred_1, y_pred_2, y_pred_3])
```

Model 1

Normality Test (Shapiro-Wilk): p-value = 0.0597

Homoscedasticity Test (Breusch-Pagan): p-value = 0.0851

Durbin-Watson Statistic: 2.1360

Model 2

Normality Test (Shapiro-Wilk): p-value = 0.0614

Homoscedasticity Test (Breusch-Pagan): p-value = 0.0855

Durbin-Watson Statistic: 2.1402

Model 3

Normality Test (Shapiro-Wilk): p-value = 0.0614

Homoscedasticity Test (Breusch-Pagan): p-value = 0.0855

Durbin-Watson Statistic: 2.1402

Con estas pruebas, además de verlo gráficamente, se puede verificar que todos los modelos cumplen los supuestos de Normalidad, Homocedasticidad e Independencia. El supuesto de

Linealidad se puede corroborar con una prueba de ANOVA. Al trabajar con Scikit-Learn, no se puede realizar la prueba de ANOVA, pero se puede verificar con la gráfica de residuos, mostrando empíricamente que el modelo es lineal.